



技术开启新“视”界
Technology Bring New Vision

基于Intel架构的实时视频流分析系统的设计与优化

翟磊 英特尔开源技术中心
2018.10.19-20 北京丽亭华苑酒店

1

背景介绍

2

Intel视觉计算平台介绍

3

Intel CS for WebRTC 实时视频流分析系统实践


4

总结



实时视频分析是智能IoT时代的重要元素

LiveVideo
StackCon
音视频技术大会



VIDEO: THE “EYE OF IOT”

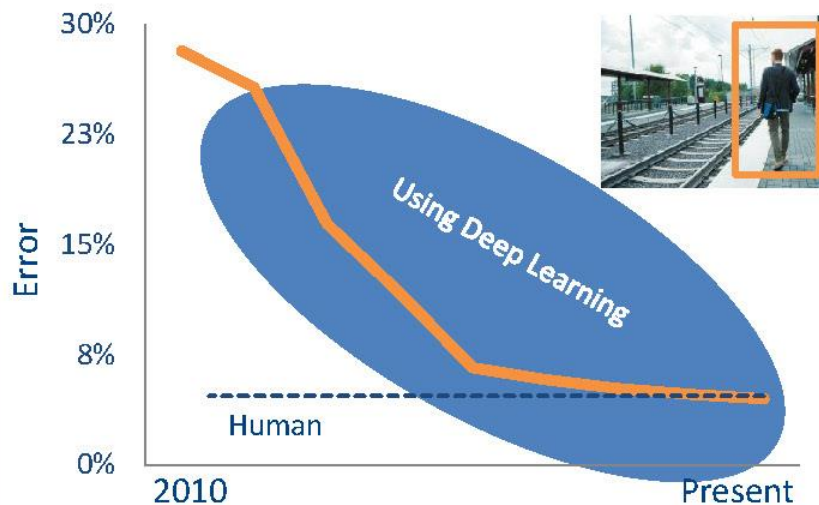
USE OF VIDEO, COMPUTER VISION, AND DEEP LEARNING IS GROWING RAPIDLY



基于深度学习的视频分析需求增长迅猛

Deep learning revenue is estimated to grow from \$655M in 2016 to **\$35B** by 2025¹

IMAGE RECOGNITION



TRADITIONAL COMPUTER VISION OBJECT DETECTION

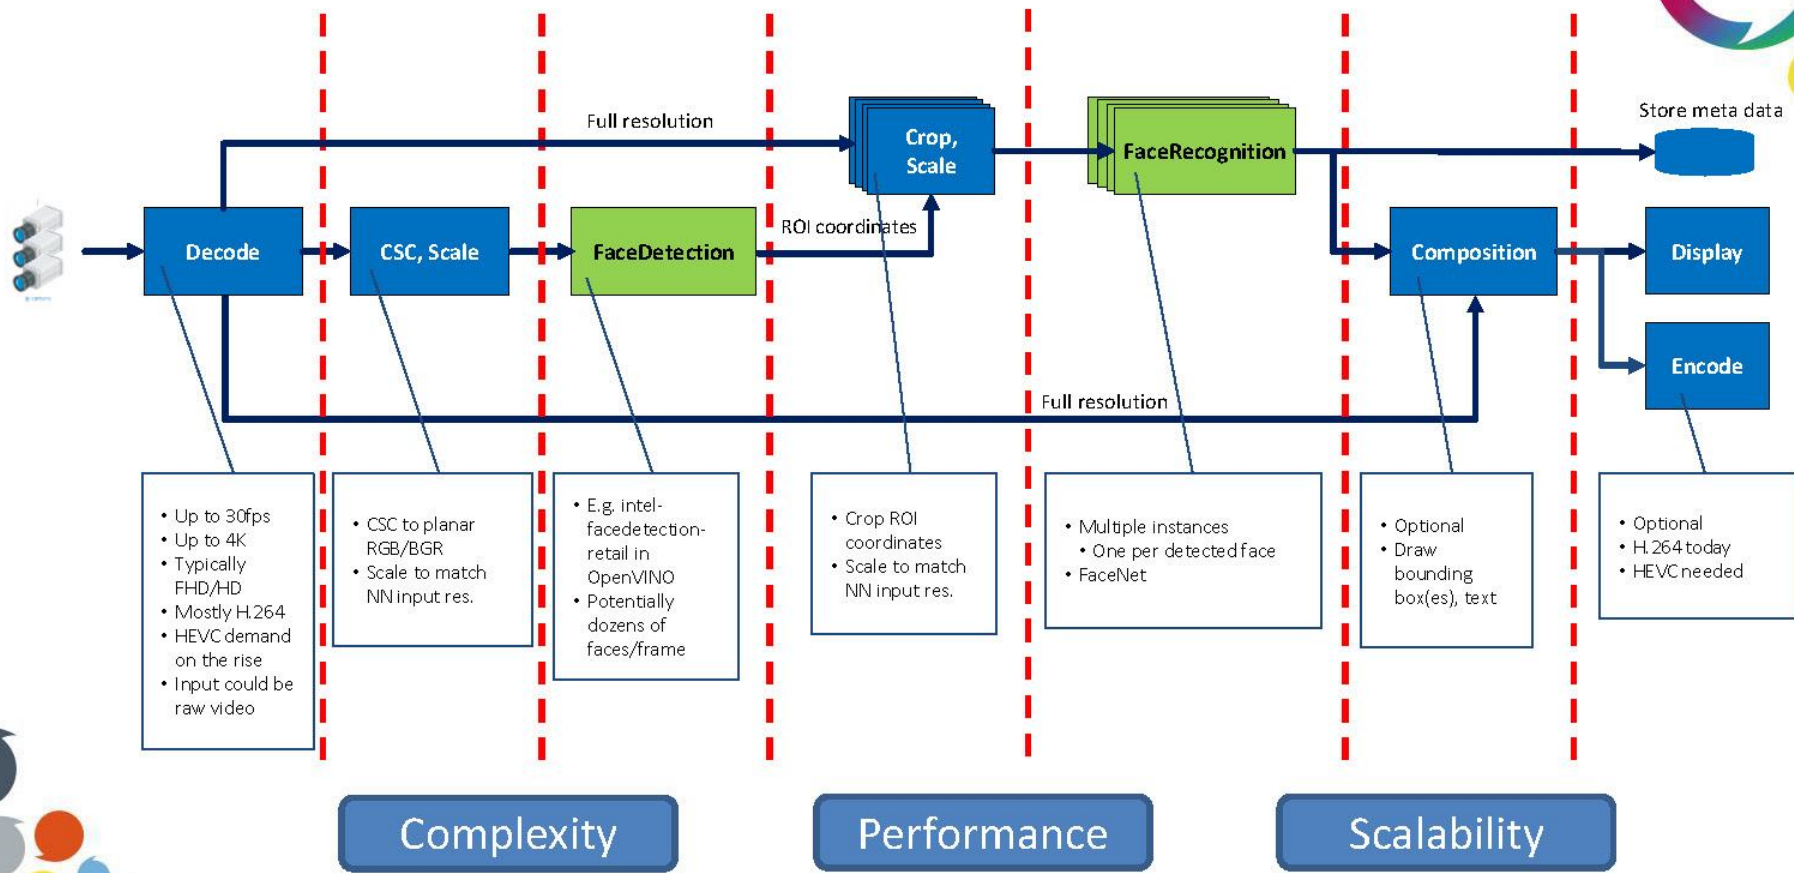


DEEP LEARNING COMPUTER VISION PERSON RECOGNITION



Market Opportunities + Advanced Technologies have Accelerated Deep Learning Adoption

实时视频流分析典型流程 (人脸识别)



Intel平台对视觉计算全方位支持

LiveVideo
StackCon
音视频技术大会

EXPERIENCES



TOOLS

Intel® Media SDK/Media Server Studio OpenVINO™ toolkit
Intel® System Studio Intel® SDK for OpenCL™ Applications
Intel® Parallel Studio XE

FRAMEWORKS



LIBRARIES

Intel® Data Analytics Acceleration Library Intel® Distribution for python Intel® Nervana™ Graph
Intel® Math Kernel Library Movidius Stack saffron TECHNOLOGY

HARDWARE



Compute



Memory & Storage



Networking



Visual Intelligence

UNLEASH FULL POTENTIAL



OpenVINO 工具包

Open Visual Inference & Neural Network Optimization



Enables deep learning on hardware accelerators and easy heterogeneous execution across Intel platforms

Intel® Deep Learning Deployment Toolkit

Model Optimizer

Convert & Optimize



Inference Engine

Optimized Inference

20+ Pre-trained Models

Computer Vision Algorithms

Samples

IR = Intermediate Representation file



Traditional Computer Vision Tools & Libraries

Optimized Libraries

OpenCV*

OpenVX*

Photography Vision

Code Samples

For Intel® CPU & CPU with integrated graphics

Increase Media/Video/Graphics Performance

Intel® Media SDK

Open Source version

OpenCL™

Drivers & Runtimes

For CPU with integrated graphics

Optimize Intel® FPGA

FPGA RunTime Environment

(from Intel® FPGA SDK for OpenCL™)

Bitstreams

FPGA – Linux* only

OS Support CentOS* 7.4 (64 bit) Ubuntu* 16.04.3 LTS (64 bit) Microsoft Windows* 10 (64 bit) Yocto Project* version Poky Jethro v2.0.3 (64 bit)

Intel® Architecture-Based
Platforms Support





Intel深度学习部署工具包

Intel Deep Learning Deployment Toolkit




Model Optimizer

- **What it is:** Preparation step -> imports trained models
- **Why important:** Optimizes for performance/space with conservative topology transformations; biggest boost is from conversion to data types matching hardware.



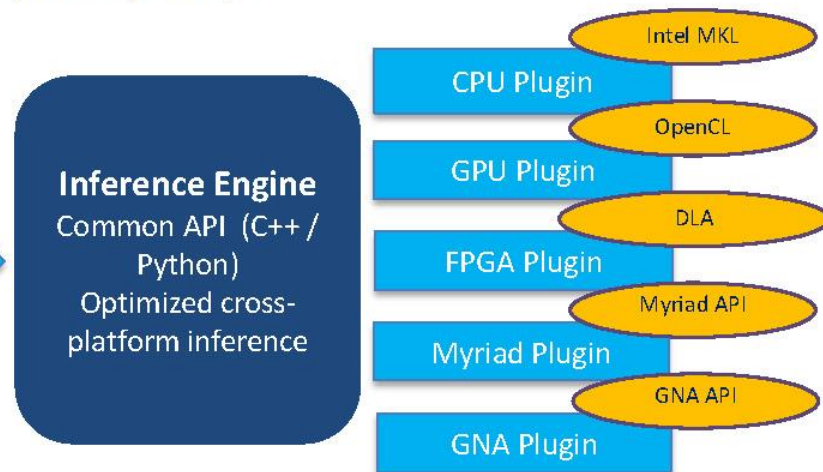
Trained Model



GPU = Intel CPU with integrated graphics processing unit/Intel® Processor Graphics

Inference Engine

- **What it is:** High-level inference API
- **Why important:** Interface is implemented as dynamically loaded plugins for each hardware type. Delivers best performance for each type without requiring users to implement and maintain multiple code pathways.





OpenVINO工具包优势



ACCELERATE PERFORMANCE

Access Intel computer vision accelerators. Speed code performance.
Supports heterogeneous processing & asynchronous execution.

INTEGRATE DEEP LEARNING

Unleash convolutional neural network (CNN) based deep learning inference using a common API, pre-trained models & computer vision algorithms.

Up to
19.9x
increase¹

SPEED DEVELOPMENT

Reduce time using a library of optimized OpenCV* & OpenVX* functions, & 15+ samples.
Develop once, deploy for current & future Intel-based devices.

INNOVATE & CUSTOMIZE

Use the increasing repository of OpenCL™ starting points in OpenCV* to add your own unique code.

Performance increase comparing certain standard framework models vs. Intel-optimized models in the Intel® Deep Learning Deployment Toolkit. Performance results are based on testing as of June 13, 2018 and may not reflect all publicly available security updates. See configuration disclosure for details. No product can be absolutely secure. For more complete information about performance and benchmark results, visit www.intel.com/benchmarks. Testing by Intel as of June 13, 2018. See Benchmark [slide 14](#) for configuration details.



Intel Media SDK & Intel Media Server Studio



Deliver fast, high quality video transcoding from camera to cloud

What it is: An API to access Intel® Quick Sync Video hardware-accelerated encode/decode & processing

- H.265 (HEVC)
- H.264 (AVC)
- MPEG-2, VP9 & more
- Resize, Scale, Deinterlace
- Color Conversion, Composition
- Denoise, Sharpen & more

Benefits

- **Boost media and video application performance** with hardware-accelerated codecs & programmable graphics on Intel® processors.**
- **Speed transition to higher frame rates & resolutions.**
- **Improve video quality, innovate cloud graphics & media analytics.**
- **Reduce infrastructure & development costs.**



FREE Downloads

software.intel.com/tools-by-segment/media ▶



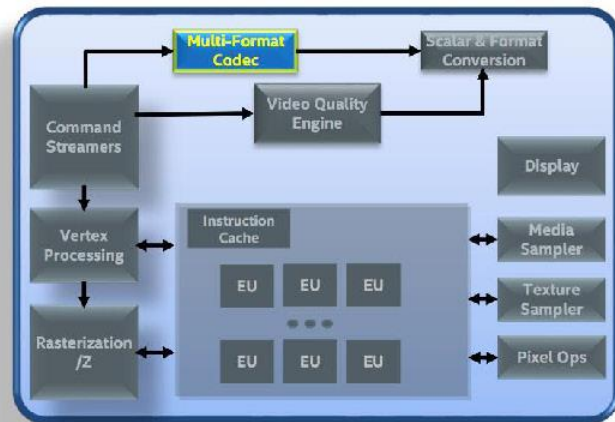
Intel Media Server Studio卓越媒体转码性能



Intel® Media Server Studio for Linux*

Multistream Performance (1xRT=30fps)		Number of Real-time (30fps) streams	Number of Real-time (60fps) streams
1080p-to-1080p	AVC-to-HEVC	15	7
	HEVC-to-HEVC	8	4
4K-to-4K	AVC-to-HEVC	4	2
	HEVC-to-HEVC	2	1

Up to 2 Real-time HEVC streams per processor¹



E3-1500 v5 HEVC is fully accelerated targeting 4K60 capability

Specific hardware technical specifications apply. See [performance benchmarks](#) and [Media Server Studio site](#) for details.

Benchmark results were obtained prior to implementation of recent software patches and firmware updates intended to address exploits referred to as "Spectre" and "Meltdown". Implementation of these updates may make these results inapplicable to your device or system. For more complete information about performance and benchmark results, visit www.intel.com/benchmarks.

Real-time HD AVC-HEVC or 4 real-time UHD AVC-HEVC transcode, 8 real-time HD HEVC-HEVC or 2 real-time UHD HEVC-HEVC transcode using Intel Media SDK (Target usage 7), all content 8-bit 4:2:0. - **Benchmark platform configuration:** Processor: Intel® Xeon® processor E3-1585Lv5 @ 3.0GHz, Ring @ 3.0GHz and GT @ 1.15GHz; primary BIOS Version: SKLSE2R1.R00.B104.B01.1511110114; driver: 20.19.15.4444, platform: RVP11 halo fab 2; OS: Windows® 8.1x64 Enterprise, 16 GB memory, 2 DIMMS 2133 MHz, one socket, four cores, Intel® Iris™ Pro Graphics P580, Intel® Hyper-threading Technology enabled, Intel® Virtualization technology enabled. Benchmark source: Intel Corporation.

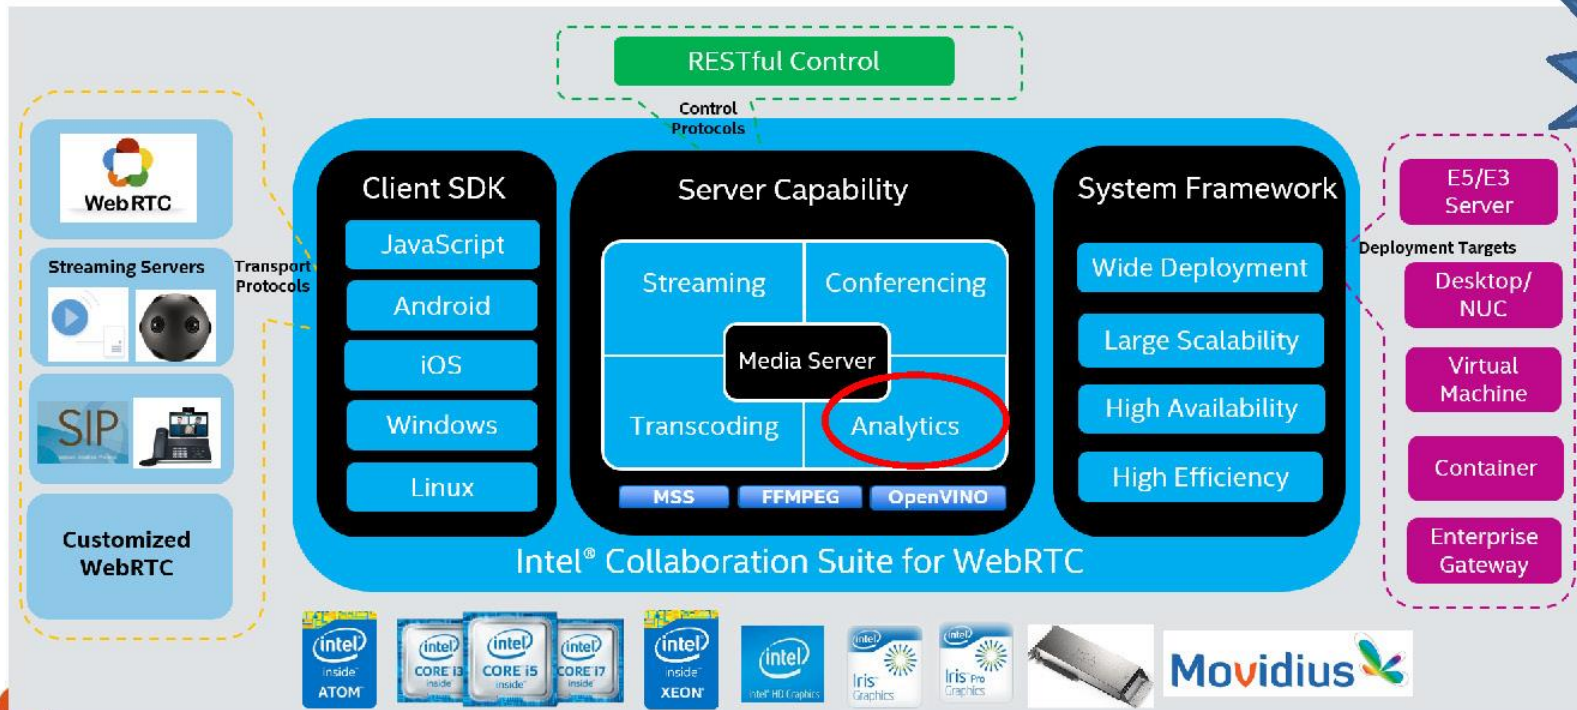
Intel's compilers may or may not optimize to the same degree for non-Intel microprocessors for optimizations that are not unique to Intel microprocessors. These optimizations include SSE2, SSE3, and SSSE3 instruction sets and other optimizations. Intel does not guarantee the availability, functionality, or effectiveness of any optimization on microprocessors not manufactured by Intel. Microprocessor-dependent optimizations in this product are intended for use with Intel microprocessors. Certain optimizations not specific to Intel microarchitecture are reserved for Intel microprocessors. Please refer to the applicable product User and Reference Guides for more information regarding the specific instruction sets covered by this notice. [Notice revision #20110804](#)

Intel Collaboration Suite for WebRTC 实时视频分析实践




An end-to-end real-time media delivery platform based on WebRTC <http://webrtc.intel.com>

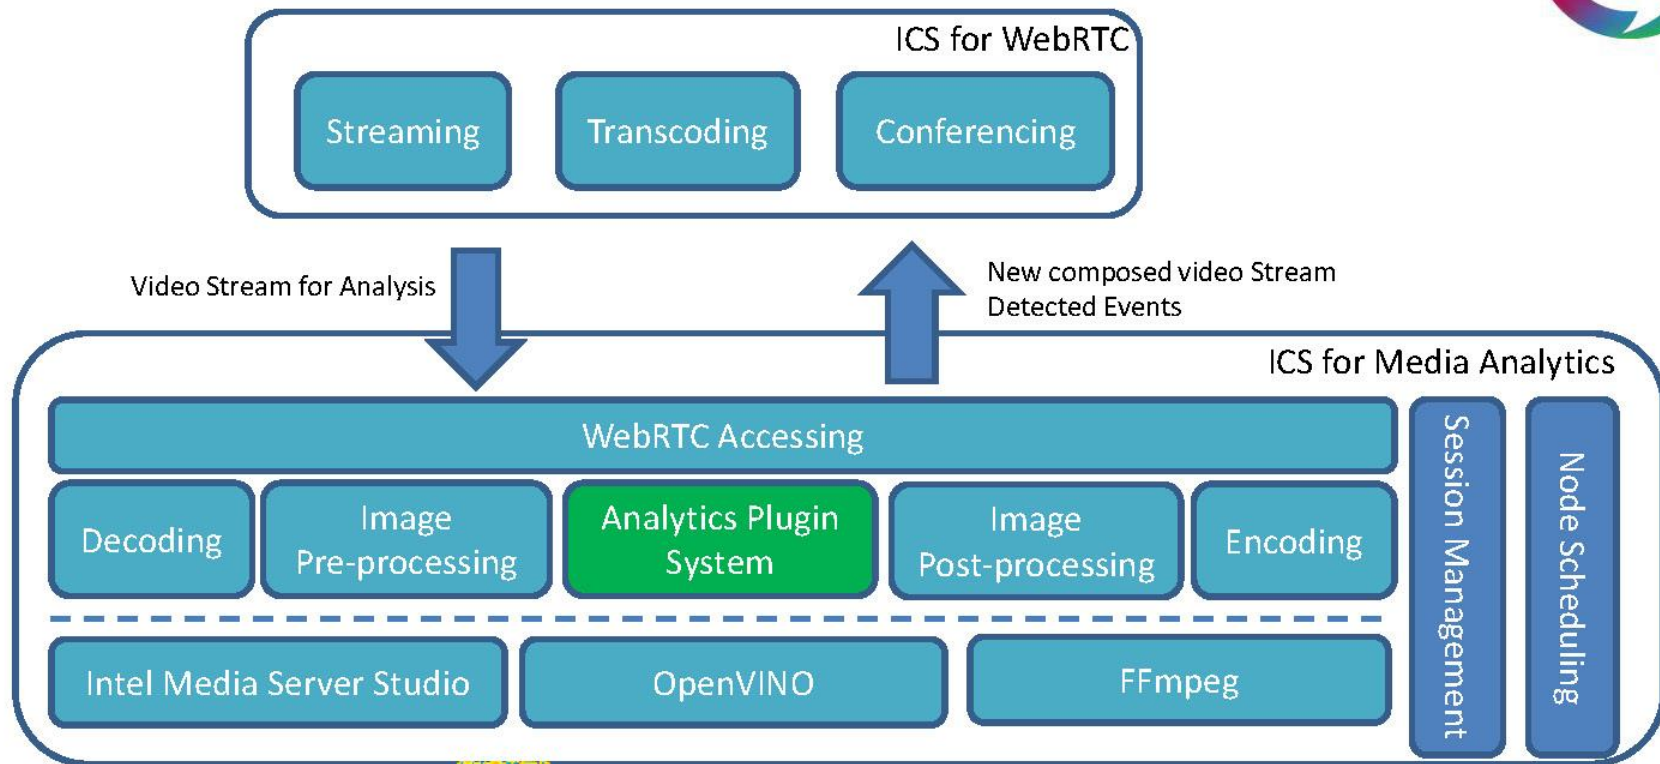
4.1
Released



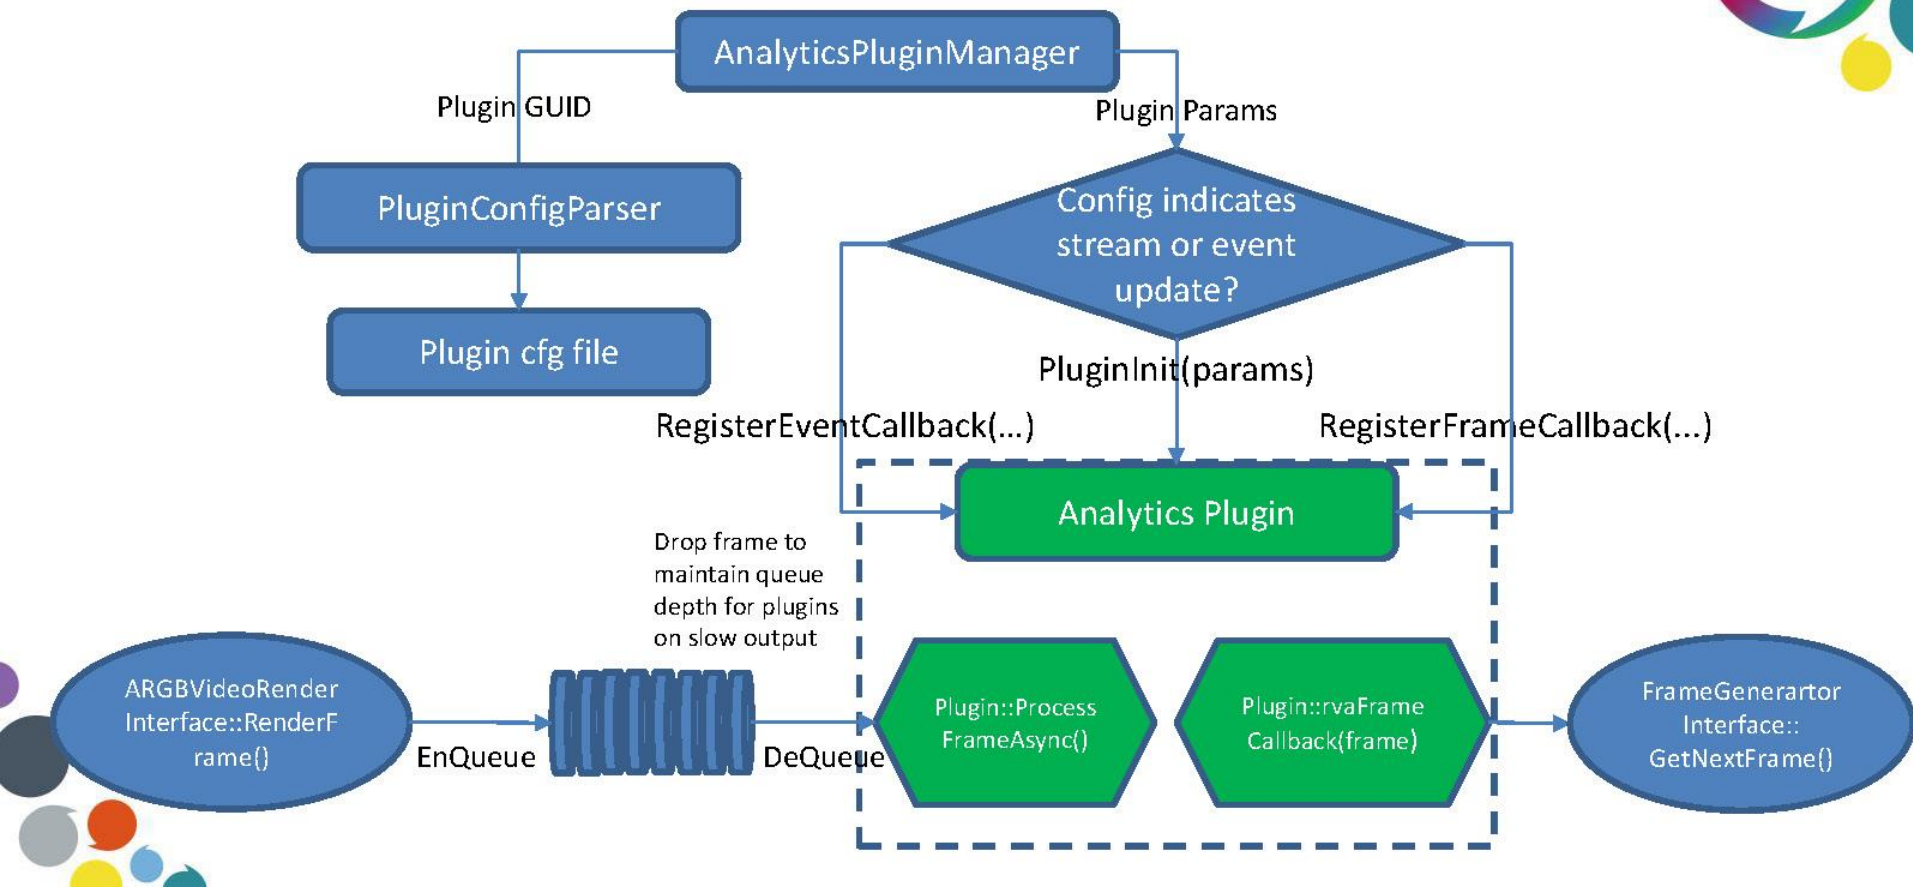
Intel CS for WebRTC 典型应用场景



Intel CS for Media Analytics 实时视频流分析系统架构



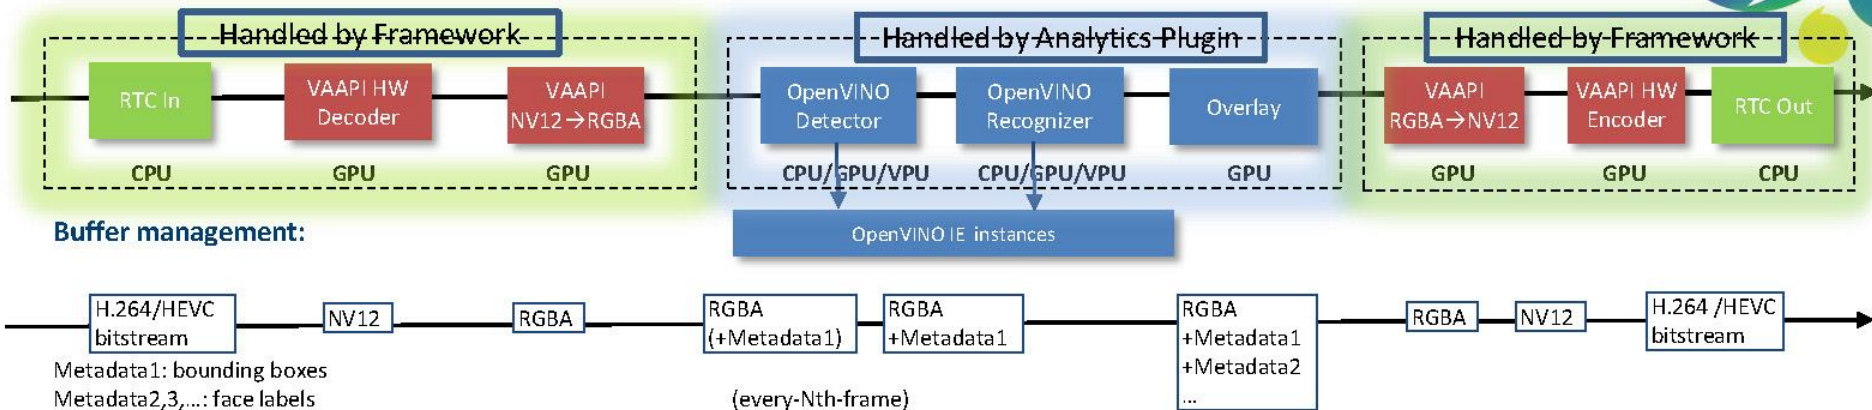
Intel CS for Media Analytics Plugin设计



Intel CS for Media Analytics分析全流程GPU加速

LiveVideo
StackCon
音视频技术大会

Single-channel source, face detection inference followed by batching face recognition inference according to detection result




Legend:

RTC components

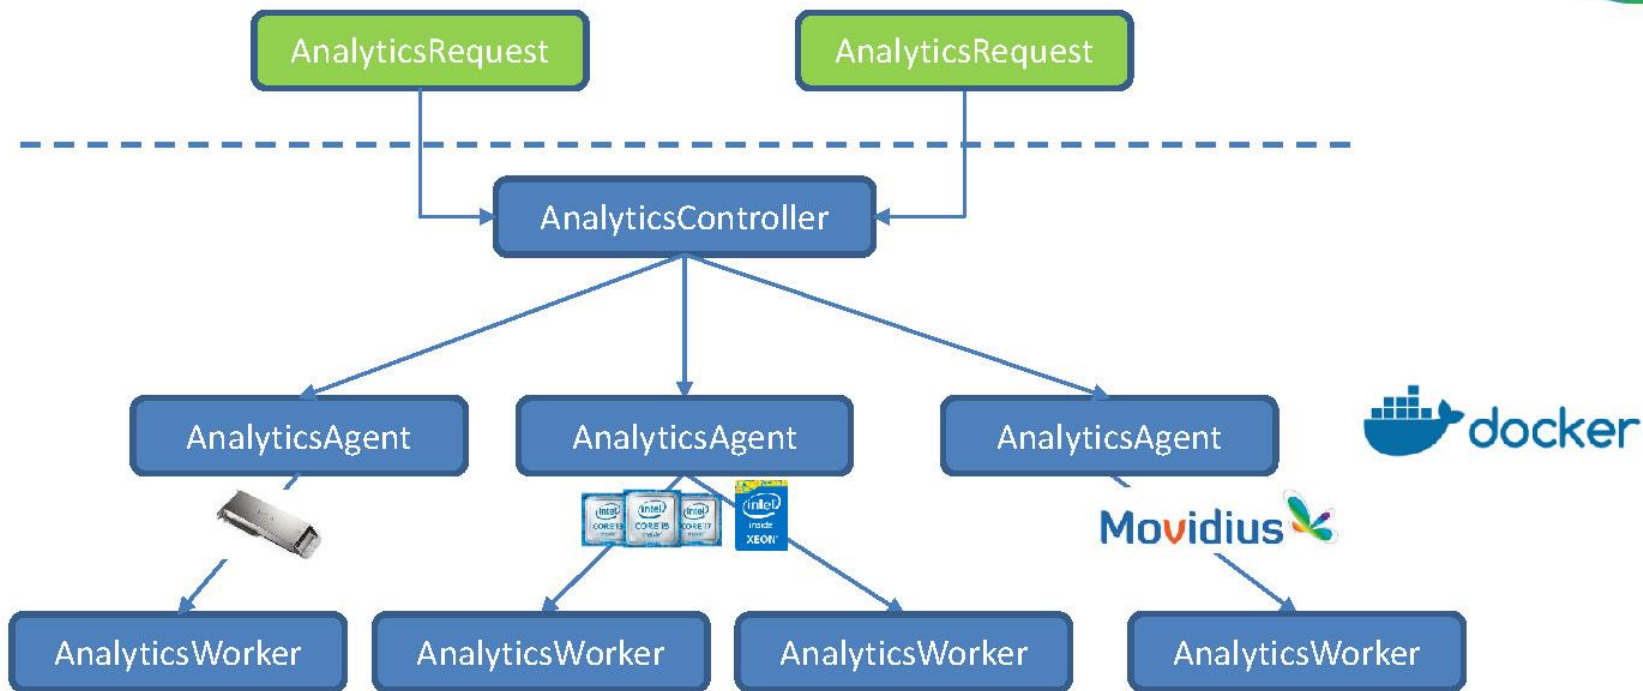
MSS components

User Analytics Plugins

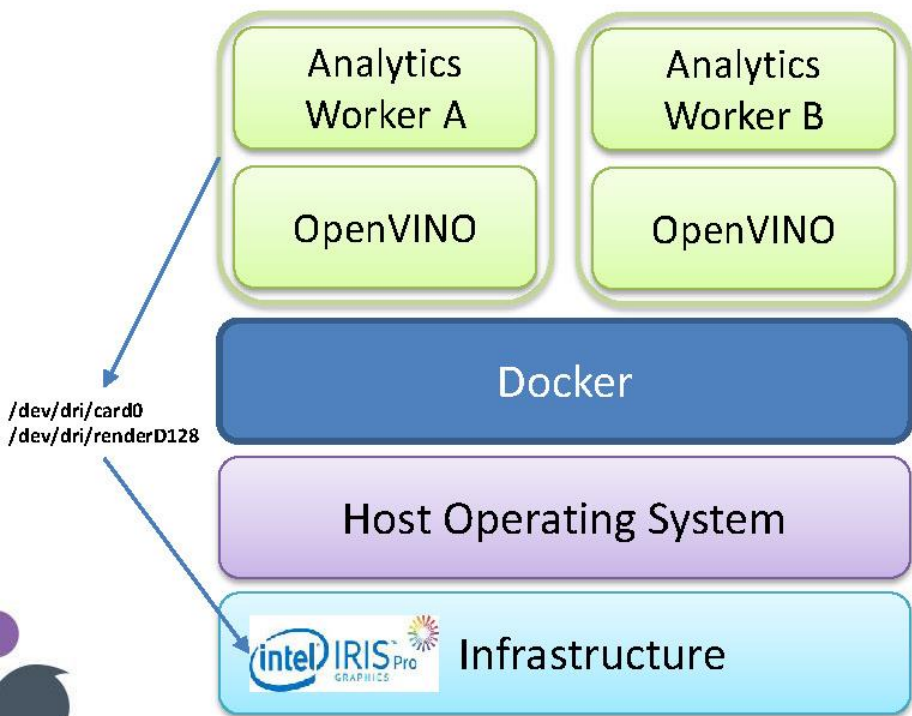
Component boundary



Intel CS for Media Analytics 系统负载优化



Analytics Agent Docker 节点Intel GPU加速实践



- Docker pull Ubuntu:16.04 image
- Run docker image w/ GPU enabled,
`--device=/dev/dri/card0 --device=/dev/dri/renderD128`
- Install dependency in container, `apt install -y cpio lsb`
- Install OpenCL™ NEO driver to enable OpenVINO GPU
- Install OpenVINO 2018R3 as **root**, see [Installation Guide](#)
- Run demo
 - `source /opt/intel/computer_vision_sdk/bin/setupvars.sh`
 - `cd /opt/intel/computer_vision_sdk/deployment_tools/demo`
 - add param "`-no_show`" to "`./security_barrier_camera_sample`" in `demo_security_barrier_camera.sh`
 - `./demo_security_barrier_camera.sh -d GPU`

Take the face recognition in classroom for example, to enable auto-call-the-roll functionality, a few challenges:

- Resource intensive due to dozens of faces
- Loss of accuracy due to scaling
- Faces at sub-picture boundaries



Intel CS for Media Analytics示例 (智能课堂点名)



Face Detection
in sub-region



Enlarge ROI on boundary and
perform detection again



Mapping
back
to original
video



Face recognition on each sub-region

Intel provided comprehensive platform and toolkit for real-time video analytics based on deep learning

Intel CS for Media Analytics + Intel Media Server Studio + OpenVINO established easy-to-use, efficient and scalable real-time video analytics framework

Intel is driving open source strategy portfolio to drive visual cloud computing industry based on IA platform

<https://www.intel.com/content/www/us/en/cloud-computing/visual-cloud.html>



Thank you

